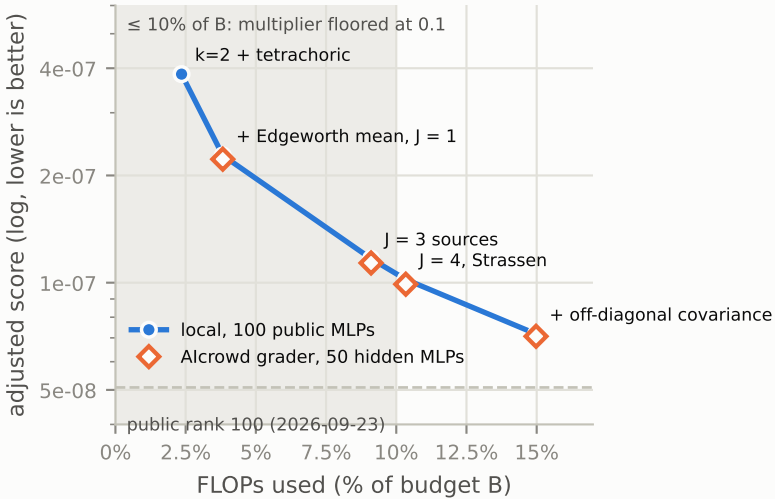
(a) every stage, scored
≤ 10% of B: multiplier floored at 0.1
k=2 + tetrachoric
+ Edgeworth mean, J = 1
J = 3 sources
J = 4, Strassen
+ off-diagonal covariance
local, 100 public MLPs
AIcrowd grader, 50 hidden MLPs
public rank 100 (2026-09-23)
adjusted score (log, lower is better)
4e-07
2e-07
1e-07
5e-08
0%
2.5%
5%
7.5%
10%
12.5%
15%
FLOPs used (% of budget B)

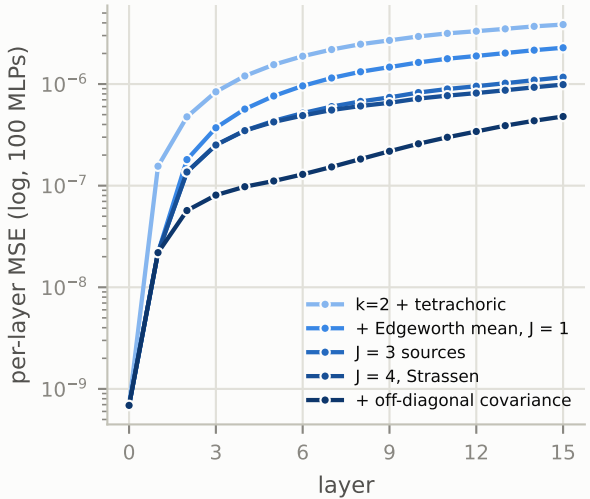
(b) the gain grows with depth
per-layer MSE (log, 100 MLPs)
10^{-6}
10^{-7}
10^{-8}
10^{-9}
k=2 + tetrachoric
+ Edgeworth mean, J = 1
J = 3 sources
J = 4, Strassen
+ off-diagonal covariance
0
3
6
9
12
15
layer